

Project: College Dataset

Stefano Cattonar, Ines El Gataa, Andrija Nicić, Angelica Rota

2025-02-10

Contents

Introduction	2
Preprocessing	2
Data Examination	2
Data Exploration	3
Observations	4
Variable visualization	6
Correlation matrix	17
Data Normalization	21
GLM	22
Dataset preparation	22
Dataset split	22
Model fitting	23
Model evaluation	24
Random Forest	24
One hundred trees model fitting	24
One hundred trees model evaluation	25
One thousand trees model fitting	28
One thousand trees model evaluation	29
Ten thousand trees model fitting	30
Ten thousand trees model evaluation	30
Comparing predictions of the GAM vs Random forest	48

```
library(dplyr)
library(corrplot)
library(caret)
```

Introduction

This project aims to analyze the **College** dataset from *An Introduction to Statistical Learning* to explore the factors that influence the number of applications received by colleges, represented by the response variable **Apps**. The dataset provides comprehensive information on various attributes of U.S. colleges, such as tuition costs, room and board expenses, financial aid, and student demographics. By examining these factors, we seek to better understand the trends and drivers behind college application rates.

Preprocessing

Data Examination

```
College = read.csv("https://www.statlearning.com/s/College.csv", header = TRUE)
help(head)
head(College)
```

```
##              X Private Apps Accept Enroll Top10perc Top25perc
## 1 Abilene Christian University    Yes 1660   1232    721      23      52
## 2      Adelphi University        Yes 2186   1924    512      16      29
## 3      Adrian College          Yes 1428   1097    336      22      50
## 4      Agnes Scott College      Yes  417    349    137      60      89
## 5 Alaska Pacific University    Yes  193    146     55      16      44
## 6      Albertson College        Yes  587    479    158      38      62
##  F.Undergrad P.Undergrad Outstate Room.Board Books Personal PhD Terminal
## 1      2885      537      7440      3300   450      2200  70      78
## 2      2683     1227     12280      6450   750      1500  29      30
## 3      1036       99     11250      3750   400      1165  53      66
## 4       510       63     12960      5450   450       875  92      97
## 5       249      869      7560      4120   800      1500  76      72
## 6       678       41     13500      3335   500       675  67      73
##  S.F.Ratio perc.alumni Expend Grad.Rate
## 1      18.1         12   7041      60
## 2      12.2         16  10527      56
## 3      12.9         30   8735      54
## 4       7.7         37  19016      59
## 5      11.9          2  10922      15
## 6       9.4         11   9727      55
```

X: name of the college

Private: A factor with levels No and Yes indicating private or public university

Accept: Number of applications accepted

Enroll: Number of new students enrolled

Top10perc: Percentage new students from top 10% of high school class

Top25perc: Percentage new students from top 25% of high school class

F.Undergrad: Number of full time undergraduates

P.Undergrad: Number of part time undergraduates

Outstate: Out-of-state tuition

Room.Board: Room and board costs

Books: Estimated book costs

Personal: Estimated personal spending

PhD: Percentage of faculty with Ph.D.'s

Terminal: Percentage of faculty with terminal degree

S.F.Ratio: Student/faculty ratio

perc.alumni: Percentage alumni who donate

Expend: Instructional expenditure per student

Grad.Rate: Graduation rate

Response variable

Apps: Number of applications received

Data Exploration

```
summary(College)
```

```
##           X              Private              Apps              Accept
## Length:777      Length:777      Min.   :   81      Min.   :   72
## Class :character Class :character 1st Qu.:  776      1st Qu.:  604
## Mode  :character Mode  :character Median : 1558      Median : 1110
##                                     Mean  : 3002      Mean   : 2019
##                                     3rd Qu.: 3624      3rd Qu.: 2424
##                                     Max.   :48094      Max.   :26330
##           Enroll          Top10perc          Top25perc          F.Undergrad
## Min.   :   35      Min.   :  1.00      Min.   :  9.0      Min.   :  139
## 1st Qu.:  242      1st Qu.:15.00      1st Qu.: 41.0      1st Qu.:  992
## Median :  434      Median :23.00      Median : 54.0      Median : 1707
## Mean   :  780      Mean   :27.56      Mean   : 55.8      Mean   : 3700
## 3rd Qu.:  902      3rd Qu.:35.00      3rd Qu.: 69.0      3rd Qu.: 4005
## Max.   :6392      Max.   :96.00      Max.   :100.0      Max.   :31643
## P.Undergrad          Outstate          Room.Board          Books
## Min.   :   1.0      Min.   : 2340      Min.   :1780      Min.   :  96.0
## 1st Qu.:  95.0      1st Qu.: 7320      1st Qu.:3597      1st Qu.: 470.0
## Median : 353.0      Median : 9990      Median :4200      Median : 500.0
## Mean   : 855.3      Mean   :10441      Mean   :4358      Mean   : 549.4
## 3rd Qu.: 967.0      3rd Qu.:12925      3rd Qu.:5050      3rd Qu.: 600.0
## Max.   :21836.0      Max.   :21700      Max.   :8124      Max.   :2340.0
## Personal            PhD              Terminal            S.F.Ratio
## Min.   :  250      Min.   :  8.00      Min.   : 24.0      Min.   :  2.50
## 1st Qu.:  850      1st Qu.: 62.00      1st Qu.: 71.0      1st Qu.:11.50
## Median :1200      Median : 75.00      Median : 82.0      Median :13.60
## Mean   :1341      Mean   : 72.66      Mean   : 79.7      Mean   :14.09
## 3rd Qu.:1700      3rd Qu.: 85.00      3rd Qu.: 92.0      3rd Qu.:16.50
## Max.   :6800      Max.   :103.00      Max.   :100.0      Max.   :39.80
## perc.alumni          Expend          Grad.Rate
```

```
## Min. : 0.00 Min. : 3186 Min. : 10.00
## 1st Qu.:13.00 1st Qu.: 6751 1st Qu.: 53.00
## Median :21.00 Median : 8377 Median : 65.00
## Mean :22.74 Mean : 9660 Mean : 65.46
## 3rd Qu.:31.00 3rd Qu.:10830 3rd Qu.: 78.00
## Max. :64.00 Max. :56233 Max. :118.00
```

Observations

1. We can note that the X column will yield no usable information.

```
College <- select(College, -X)
```

2. We have to exclude also Accept and Enroll to prevent data leakage.

```
College <- select(College, -c(Accept, Enroll))
summary(College)
```

```
## Private Apps Top10perc Top25perc
## Length:777 Min. : 81 Min. : 1.00 Min. : 9.0
## Class :character 1st Qu.: 776 1st Qu.:15.00 1st Qu.: 41.0
## Mode :character Median : 1558 Median :23.00 Median : 54.0
## Mean : 3002 Mean :27.56 Mean : 55.8
## 3rd Qu.: 3624 3rd Qu.:35.00 3rd Qu.: 69.0
## Max. :48094 Max. :96.00 Max. :100.0
## F.Undergrad P.Undergrad Outstate Room.Board
## Min. : 139 Min. : 1.0 Min. : 2340 Min. :1780
## 1st Qu.: 992 1st Qu.: 95.0 1st Qu.: 7320 1st Qu.:3597
## Median : 1707 Median : 353.0 Median : 9990 Median :4200
## Mean : 3700 Mean : 855.3 Mean :10441 Mean :4358
## 3rd Qu.: 4005 3rd Qu.: 967.0 3rd Qu.:12925 3rd Qu.:5050
## Max. :31643 Max. :21836.0 Max. :21700 Max. :8124
## Books Personal PhD Terminal
## Min. : 96.0 Min. : 250 Min. : 8.00 Min. : 24.0
## 1st Qu.: 470.0 1st Qu.: 850 1st Qu.: 62.00 1st Qu.: 71.0
## Median : 500.0 Median :1200 Median : 75.00 Median : 82.0
## Mean : 549.4 Mean :1341 Mean : 72.66 Mean : 79.7
## 3rd Qu.: 600.0 3rd Qu.:1700 3rd Qu.: 85.00 3rd Qu.: 92.0
## Max. :2340.0 Max. :6800 Max. :103.00 Max. :100.0
## S.F.Ratio perc.alumni Expend Grad.Rate
## Min. : 2.50 Min. : 0.00 Min. : 3186 Min. : 10.00
## 1st Qu.:11.50 1st Qu.:13.00 1st Qu.: 6751 1st Qu.: 53.00
## Median :13.60 Median :21.00 Median : 8377 Median : 65.00
## Mean :14.09 Mean :22.74 Mean : 9660 Mean : 65.46
## 3rd Qu.:16.50 3rd Qu.:31.00 3rd Qu.:10830 3rd Qu.: 78.00
## Max. :39.80 Max. :64.00 Max. :56233 Max. :118.00
```

3. All the variables does not have negative values.
4. The variables does not have NA, so we don't need to replace the missing values

5. We want to check if we have some duplicates

```
idx <- which(duplicated(College))  
idx
```

```
## integer(0)
```

6. The variable `Private` is a categorical one so we convert its values.

```
College$Private <- as.factor(College$Private)  
College$Private <- as.numeric(College$Private) - 1 # "Yes" is 1, "No" is 0  
College$Private[11:30]
```

```
## [1] 1 1 1 1 1 1 1 1 1 0 1 0 1 0 1 0 1 0 1 1
```

7. We can note that `Private` is not balanced

```
index_dataset_public <- which(College$Private == 0)  
index_dataset_private <- which(College$Private == 1)  
private_proportion <- (length(index_dataset_private) / length(index_dataset_public))  
table(College$Private)
```

```
##  
## 0 1  
## 212 565
```


```
private_proportion
```

```
## [1] 2.665094
```

We can see from `private_proportion` that “Yes” is 2.66 times more frequent than “No”.

Now we can see the barplot

```
barplot(table(College$Private), main="Private")
```


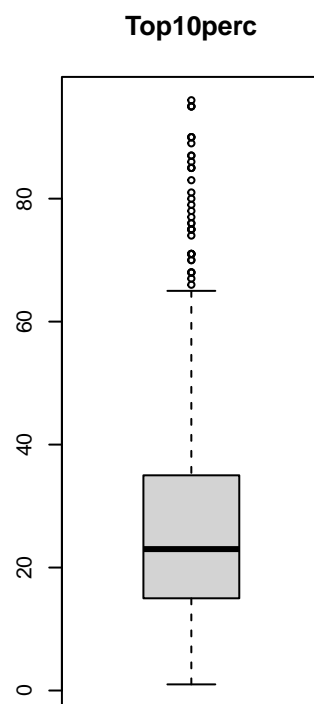
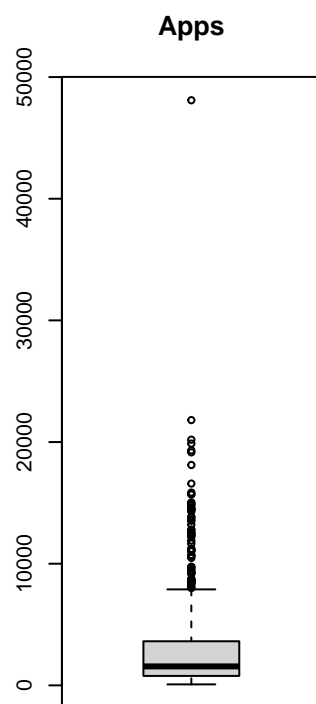





Variable visualization



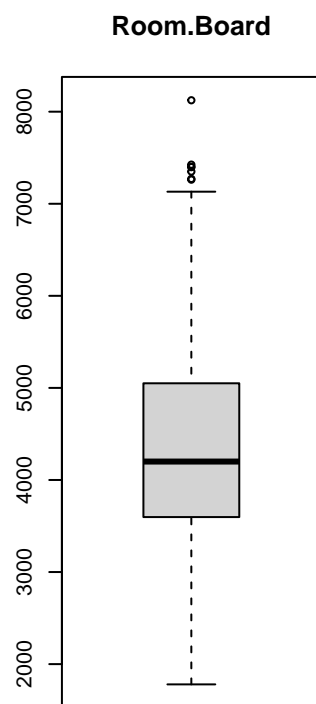
Now we can visualize the variables and their distributions.


Boxplot, excluded `Private` that is a categorical variable


```
par(mfrow=c(1,3))
for (i in 1:length(College)) {
  if ( names(College[i]) != "Private"){
    boxplot(College[,i], main=names(College[i]))
  }
}
```























Histogram, excluded `Private` that is a categorical variable

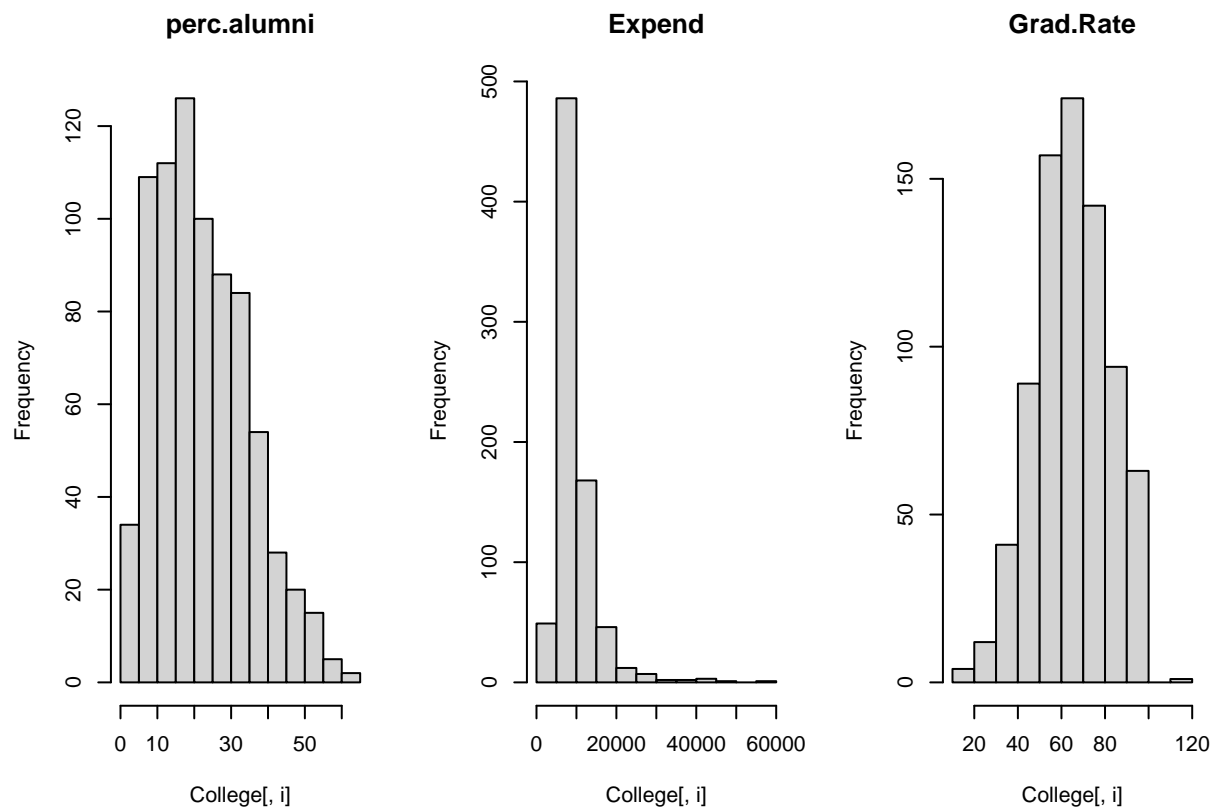
```
par(mfrow=c(1,3))
for (i in 1:length(College)) {
  if ( names(College[i]) != "Private"){
    hist(College[,i], main=names(College[i]))
  }
}
```











We can see that F.Undergrad and P.Undergrad are really skewed so we can take the logarithm.


```
College$F.Undergrad <- log(College$F.Undergrad)
College$P.Undergrad <- log(College$P.Undergrad)

par(mfrow = c(2, 2))


boxplot(College$F.Undergrad, main="Full time undergraduates")
boxplot(College$P.Undergrad, main="Part time undergraduates")

hist(College$F.Undergrad, main="Full time undergraduates")
hist(College$P.Undergrad, main="Part time undergraduates")
```

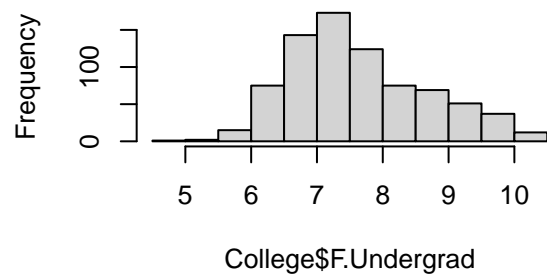

Full time undergraduates



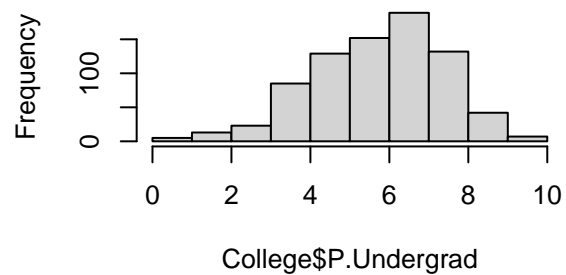
Part time undergraduates



Full time undergraduates



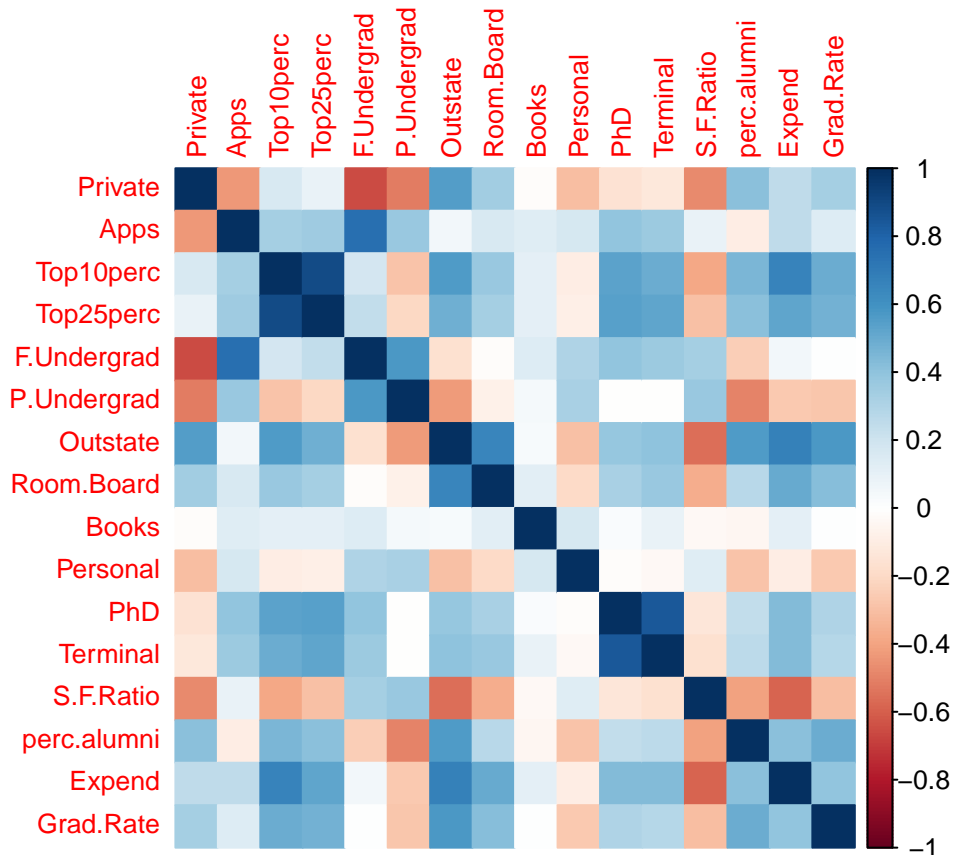
Part time undergraduates



Correlation matrix

We plot the correlation matrix to examine the relationships between the predictors and Apps

```
correlation_matrix <- cor(College)
corrplot(correlation_matrix, method="color", tl.cex = 0.8)
```

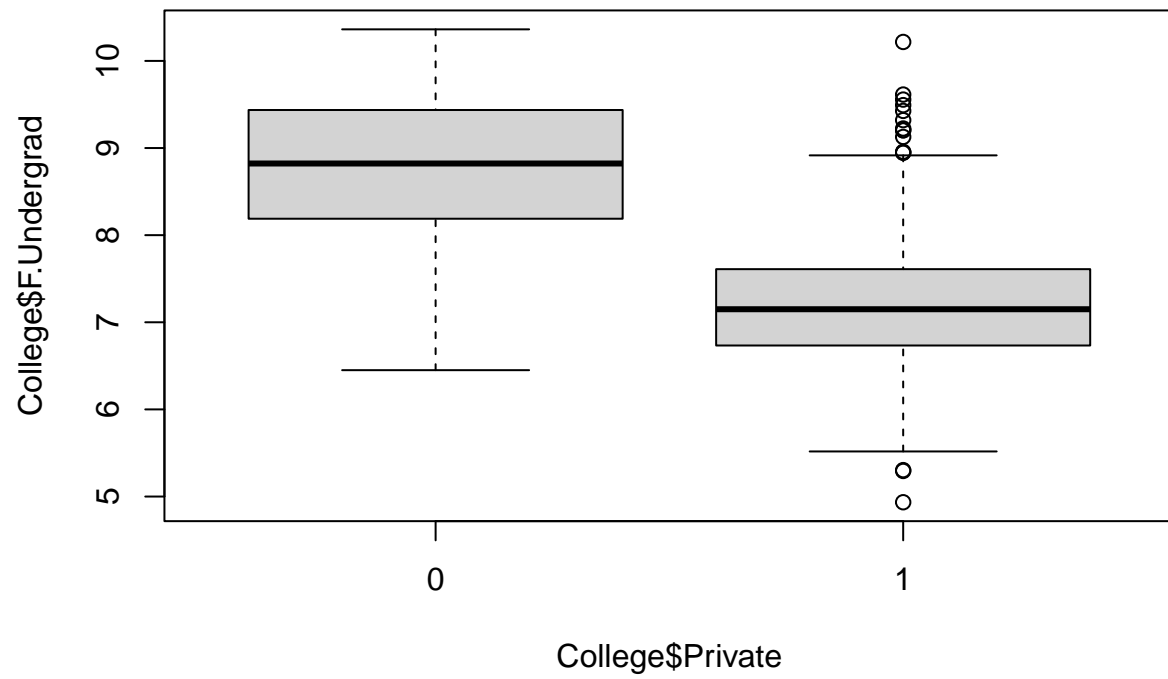


As you can see from the correlation matrix, there are some variables that are correlated with each other. We can see that the variables **Terminal** and **PhD** are really directly heavily correlated. This phenomena is caused by the fact that **PhD** contains a sub group of **Terminal**. So we can exclude **PhD**.

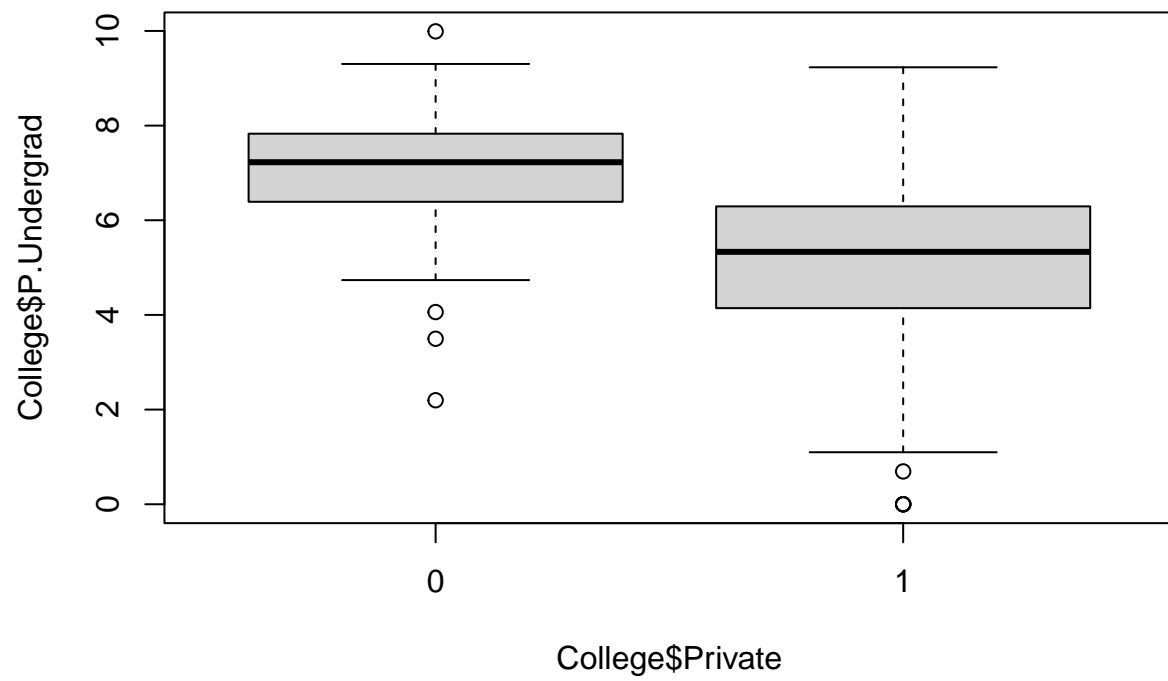
Top10perc and **Top25perc** are the same so for the same reason we can exclude **Top10perc**.

Now we can visualize some of the correlated variables

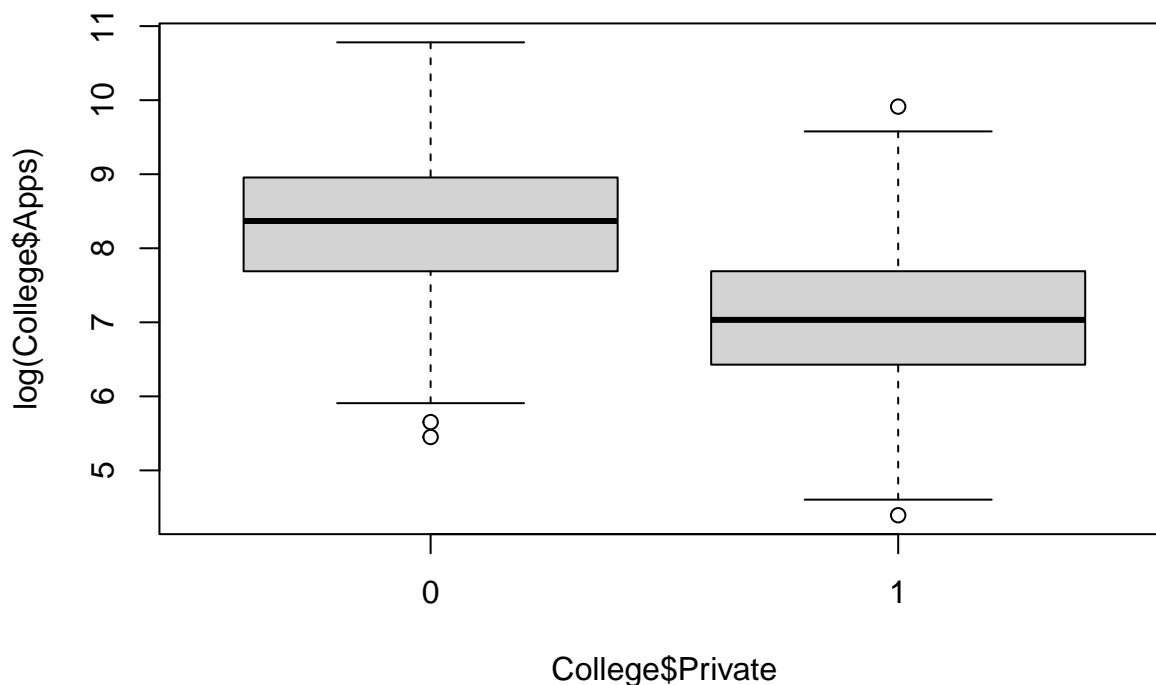
```
boxplot(College$F.Undergrad ~ College$Private)
```



```
boxplot(College$P.Undergrad ~ College$Private)
```



```
boxplot(log(College$Apps) ~ College$Private)
```



Data Normalization

Now let's normalize the data so that we have all the data with comparable scales.

```
normalize <- function(x) {
  return((x - min(x)) / (max(x) - min(x)))
}

college_norm <- College
college_norm[, 2:ncol(College)] <- lapply(College[, 2:ncol(College)], normalize)

summary(college_norm)
```

```
##      Private      Apps      Top10perc      Top25perc
## Min.   :0.0000   Min.   :0.00000   Min.   :0.0000   Min.   :0.0000
## 1st Qu.:0.0000   1st Qu.:0.01448   1st Qu.:0.1474   1st Qu.:0.3516
## Median :1.0000   Median :0.03076   Median :0.2316   Median :0.4945
## Mean   :0.7272   Mean   :0.06083   Mean   :0.2796   Mean   :0.5142
## 3rd Qu.:1.0000   3rd Qu.:0.07379   3rd Qu.:0.3579   3rd Qu.:0.6593
## Max.   :1.0000   Max.   :1.00000   Max.   :1.0000   Max.   :1.0000
## F.Undergrad P.Undergrad Outstate Room.Board
## Min.   :0.0000   Min.   :0.0000   Min.   :0.0000   Min.   :0.0000
## 1st Qu.:0.3621   1st Qu.:0.4558   1st Qu.:0.2572   1st Qu.:0.2864
## Median :0.4621   Median :0.5872   Median :0.3951   Median :0.3815
```

## Mean	:0.4976	Mean	:0.5696	Mean	:0.4184	Mean	:0.4063
## 3rd Qu.	:0.6192	3rd Qu.	:0.6880	3rd Qu.	:0.5467	3rd Qu.	:0.5154
## Max.	:1.0000	Max.	:1.0000	Max.	:1.0000	Max.	:1.0000
## Books		Personal		PhD		Terminal	
## Min.	:0.0000	Min.	:0.0000	Min.	:0.0000	Min.	:0.0000
## 1st Qu.	:0.1667	1st Qu.	:0.0916	1st Qu.	:0.5684	1st Qu.	:0.6184
## Median	:0.1800	Median	:0.1450	Median	:0.7053	Median	:0.7632
## Mean	:0.2020	Mean	:0.1665	Mean	:0.6806	Mean	:0.7329
## 3rd Qu.	:0.2246	3rd Qu.	:0.2214	3rd Qu.	:0.8105	3rd Qu.	:0.8947
## Max.	:1.0000	Max.	:1.0000	Max.	:1.0000	Max.	:1.0000
## S.F.Ratio		perc.alumni		Expend		Grad.Rate	
## Min.	:0.0000	Min.	:0.0000	Min.	:0.00000	Min.	:0.0000
## 1st Qu.	:0.2413	1st Qu.	:0.2031	1st Qu.	:0.06720	1st Qu.	:0.3981
## Median	:0.2976	Median	:0.3281	Median	:0.09786	Median	:0.5093
## Mean	:0.3107	Mean	:0.3554	Mean	:0.12205	Mean	:0.5135
## 3rd Qu.	:0.3753	3rd Qu.	:0.4844	3rd Qu.	:0.14410	3rd Qu.	:0.6296
## Max.	:1.0000	Max.	:1.0000	Max.	:1.00000	Max.	:1.0000

GLM

Dataset preparation

Now we start to clean the dataset as told before:

```
#Copy the dataset
college_cleared <- college_norm

#Remove the target variable and the variables that are not significant
college_cleared$Accept <- NULL
college_cleared$Enroll <- NULL
college_cleared$X <- NULL
college_cleared$Apps <- NULL
college_cleared$Top10perc <- NULL
college_cleared$PhD <- NULL
```

Dataset split

Now we must split the dataset in train set and test set. We will use 80% of the dataset for the training and 20% for the test.

```
# Set a seed to make it deterministic
set.seed(42)

# Calculate the number of training lines (80% of the dataset)
n_training_lines <- floor(0.8 * nrow(college_cleared))

# Ensure the number of training lines is even
if (n_training_lines %% 2 != 0) {
  n_training_lines <- n_training_lines - 1
}
```

```

# save index of lines with Private 1 and lines with Private 0
index_dataset_public <- which(college_cleared$Private == 0)

index_dataset_private <- which(college_cleared$Private == 1)

# Calculate the exact percentage of private 1 lines in the dataset
percentage_private <- length(index_dataset_private) / nrow(college_cleared)

# take 80% of lines from index_dataset_public
train_index_dataset_public <- sample(index_dataset_public, floor(n_training_lines * (1 - percentage_private)))
train_index_dataset_private <- sample(index_dataset_private, floor(n_training_lines * (percentage_private)))

# Cobine the two train_index_dataset_* into train_index
train_index <- c(train_index_dataset_public, train_index_dataset_private)

train_data <- college_cleared[train_index,]
test_data <- college_cleared[-train_index,]

```

Model fitting

Now we can fit the glm on the training set:

```
glm <- glm(College$Apps[train_index] ~ ., data=train_data, family=poisson)
```

```
summary(glm)
```

```

##
## Call:
## glm(formula = College$Apps[train_index] ~ ., family = poisson,
##      data = train_data)
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  4.092305   0.008438  485.009 < 2e-16 ***
## Private      -0.184864   0.002964  -62.360 < 2e-16 ***
## Top25perc     0.177462   0.004820   36.818 < 2e-16 ***
## F.Undergrad   4.729292   0.007555  625.964 < 2e-16 ***
## P.Undergrad  -0.146204   0.007927  -18.443 < 2e-16 ***
## Outstate      0.492628   0.007140   68.997 < 2e-16 ***
## Room.Board    0.748332   0.006285  119.061 < 2e-16 ***
## Books          0.282714   0.010330   27.368 < 2e-16 ***
## Personal     -0.350038   0.008107  -43.179 < 2e-16 ***
## Terminal       0.019375   0.006914    2.802  0.00508 **
## S.F.Ratio      0.351358   0.009843   35.696 < 2e-16 ***
## perc.alumni  -0.049954   0.005894   -8.475 < 2e-16 ***
## Expend         0.570555   0.010381   54.959 < 2e-16 ***
## Grad.Rate      0.955605   0.006981  136.879 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for poisson family taken to be 1)
##

```

```
## Null deviance: 2120079 on 618 degrees of freedom
## Residual deviance: 221113 on 605 degrees of freedom
## AIC: 226872
##
## Number of Fisher Scoring iterations: 4
```

Model evaluation

Now we can calculate the RMSE for the glm:

```
RMSE <- function(predicted, actual) {
  sqrt(mean((predicted - actual)^2))
}

rmse_train_set <- RMSE(predict.glm(glm), College$Apps[train_index])
rmse_test_set <- RMSE(predict.glm(glm, newdata = test_data), College$Apps[-train_index])
```

The rmse for the train_set is inside rmse_train_set:

```
rmse_train_set
```

```
## [1] 5009.428
```

The rmse for the test_set is inside rmse_test_set:

```
rmse_test_set
```

```
## [1] 4394.249
```

The RMSE for the test set is lower than the RMSE for the train set, which is a good sign that the model is not overfitting and it is generalizing well.

Random Forest

One hundred trees model fitting

Now, we will try to fit a random forest model on the training set and see if it performs better than the glm:

```
library(randomForest)
```

```
## randomForest 4.7-1.2
```

```
## Type rfNews() to see new features/changes/bug fixes.
```

```
##
```

```
## Attaching package: 'randomForest'
```



```
## The following object is masked from 'package:ggplot2':
##
##   margin
```

```
## The following object is masked from 'package:dplyr':
##
##   combine
```

```
library(randomForestExplainer)
```

```
## Registered S3 method overwritten by 'GGally':
##   method from
##   +.gg      ggplot2
```

```
random_forest <- randomForest(x=train_data, formula=Apps ~ ., y=College$Apps[train_index], data = train,
                               ntree=100, mtry=4, importance=TRUE)
random_forest
```

```
##
## Call:
## randomForest(x = train_data, y = College$Apps[train_index], xtest = test_data, ytest = College$Apps[test_index],
##              type = "regression", ntree = 100, mtry = 4, importance = TRUE)
##              Type of random forest: regression
##              Number of trees: 100
## No. of variables tried at each split: 4
##
##              Mean of squared residuals: 3851069
##              % Var explained: 76.05
##              Test set MSE: 1992994
##              % Var explained: 81.22
```

One hundred trees model evaluation

Now we can calculate the RMSE for the random forest model:

```
#RMSE train_set
rf_rmse_train_set <- sqrt(random_forest$mse[length(random_forest$mse)])

#RMSE test_set
rf_rmse_test_set <- sqrt(random_forest$test$mse[length(random_forest$test$mse)])
```

The rmse for the train_set is inside rf_rmse_train_set:

```
rf_rmse_train_set
```

```
## [1] 1962.414
```

The rmse for the test_set is inside rf_rmse_test_set:

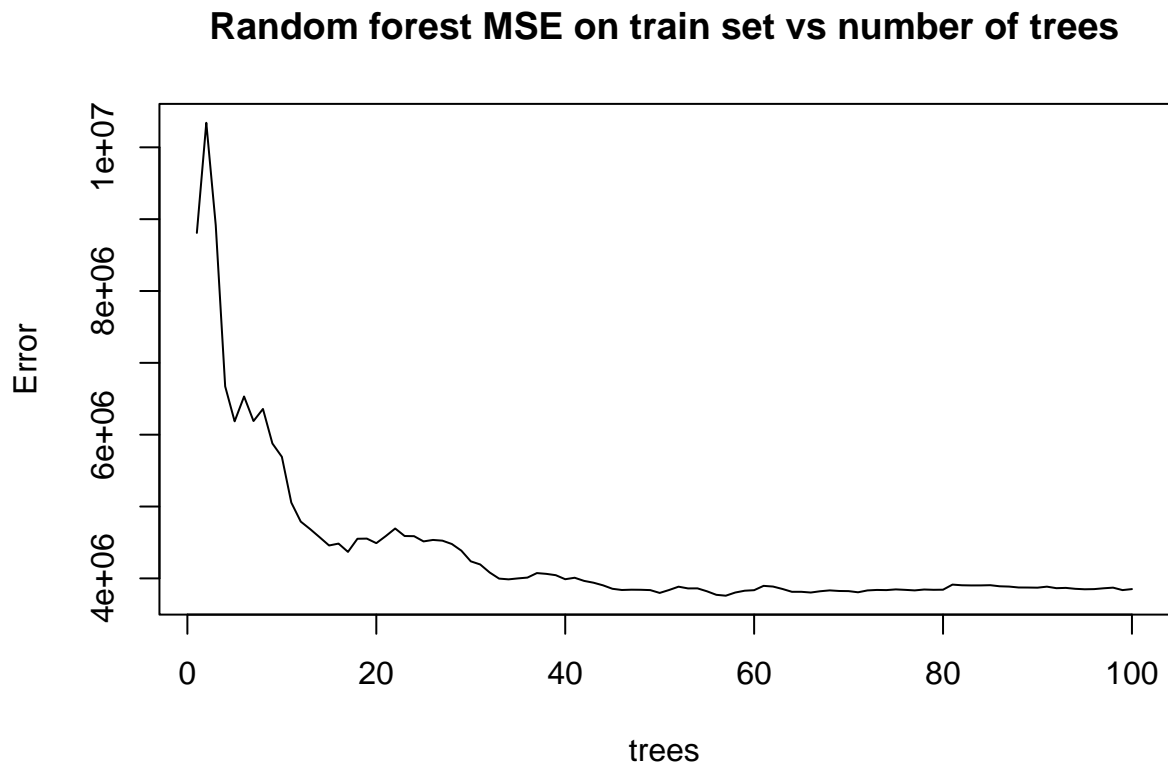
```
rf_rmse_test_set
```

```
## [1] 1411.734
```

The RMSE of the random forest model is lower than the RMSE of the glm model, which means that the random forest model is performing better than the glm model.

Now we can plot how the MSE of the model changes with the number of trees:

```
plot(random_forest, main = "Random forest MSE on train set vs number of trees")
```



Now explore how the variables are used in the random forest:

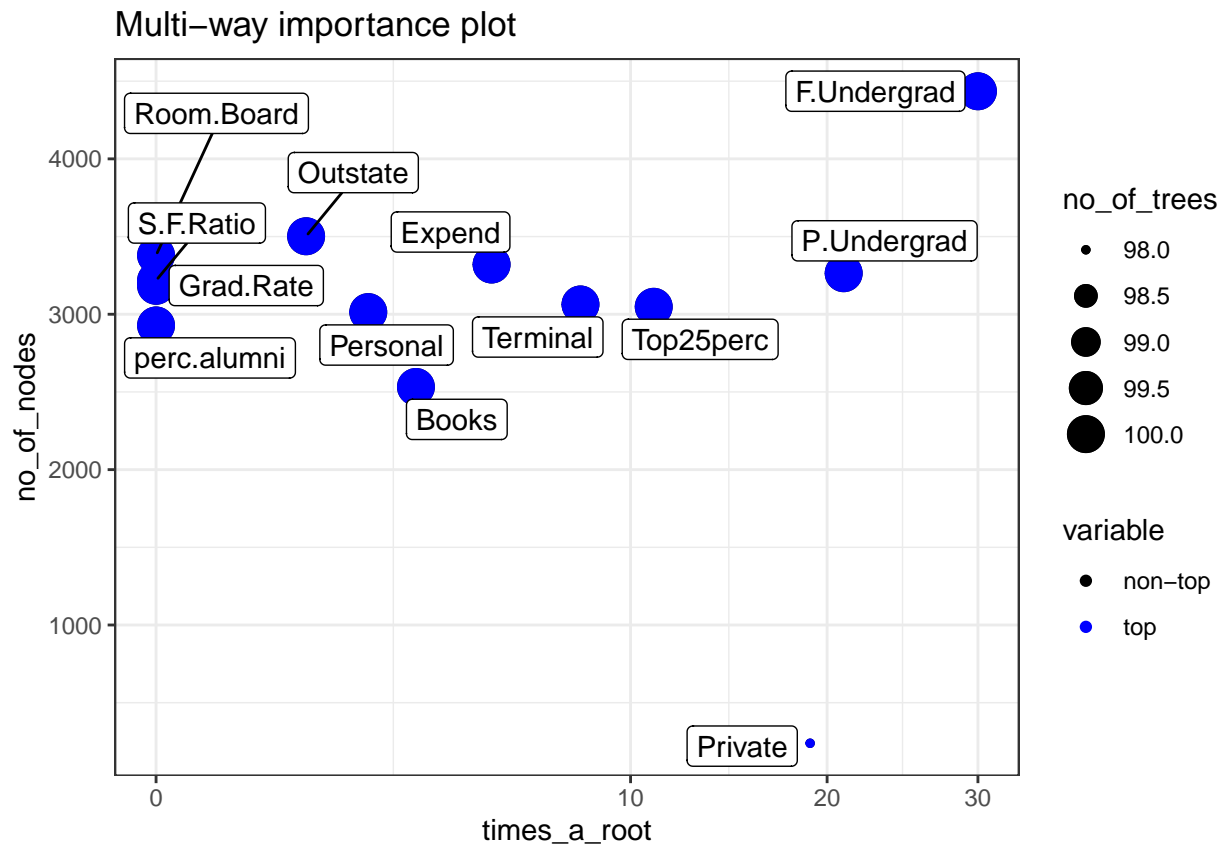
```
importance_random_forest <- measure_importance(random_forest)
```

```
importance_random_forest
```

```
##      variable mean_min_depth no_of_nodes mse_increase node_purity_increase
## 1      Books          3.7400         2532    -270617.3         234728137
## 2     Expend          2.3900         3320     968186.6         595865266
## 3 F.Undergrad          1.1600         4435   13947069.3         3853122354
## 4   Grad.Rate          2.7700         3183    2870676.7         620419170
## 5   Outstate          2.6800         3502    1364694.6         494687655
## 6 P.Undergrad          2.1800         3264    2763740.6         1111092839
## 7 perc.alumni          3.9300         2928     147399.1         189797185
```

## 8	Personal	3.8400	3014	186349.0	202430936
## 9	Private	4.2166	239	3283327.3	562522412
## 10	Room.Board	3.2100	3379	460229.4	283365283
## 11	S.F.Ratio	3.2400	3217	151082.2	332505229
## 12	Terminal	2.6700	3063	416483.8	498339608
## 13	Top25perc	2.1400	3048	960212.7	944531722
##	no_of_trees	times_a_root	p_value		
## 1	100	3	1.000000e+00		
## 2	100	5	3.420597e-09		
## 3	100	30	1.329068e-142		
## 4	100	0	5.648074e-04		
## 5	100	1	3.814932e-20		
## 6	100	21	9.631230e-07		
## 7	100	0	9.406634e-01		
## 8	100	2	4.689851e-01		
## 9	98	19	1.000000e+00		
## 10	100	0	3.031443e-12		
## 11	100	0	5.028945e-05		
## 12	100	8	1.574935e-01		
## 13	100	11	2.352934e-01		

```
plot_multi_way_importance(importance_random_forest, y_measure = "no_of_nodes", x_measure = "times_a_root"
```



One thousand trees model fitting

And if we increase the number of trees from 100 to 1000:

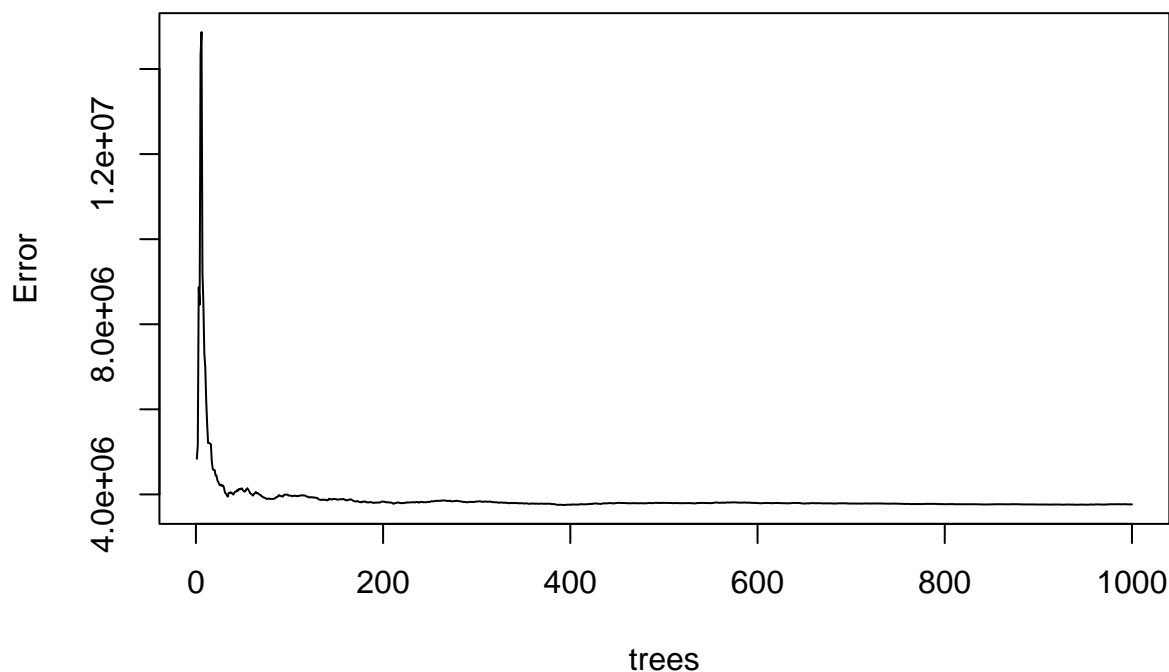
```
random_forest_1000 <- randomForest(x=train_data, formula=Apps ~ ., y=College$Apps[train_index], data = train_data)

random_forest_1000

##
## Call:
## randomForest(x = train_data, y = College$Apps[train_index], xtest = test_data, ytest = College$Apps[test_index],
##               type = "regression", number = 1000, variables.tried = 4,
##               mean_squared_residual = 3766066, var_explained = 76.58, test_mse = 1739743, var_explained_test = 83.61)
## No. of variables tried at each split: 4
##
##               Mean of squared residuals: 3766066
##               % Var explained: 76.58
##               Test set MSE: 1739743
##               % Var explained: 83.61

plot(random_forest_1000, main = "Random forest MSE on train set vs number of trees")
```

Random forest MSE on train set vs number of trees



```
#RMSE train_set
rf_rmse_train_set_1000 <- sqrt(random_forest_1000$mse[length(random_forest_1000$mse)])
```

```
#RMSE test_set
rf_rmse_test_set_1000 <- sqrt(random_forest_1000$test$mse[length(random_forest_1000$test$mse)])
```

One thousand trees model evaluation

Analyze Root Mean Square Error for the random forest model with 1000 trees:

```
rf_rmse_train_set_1000
```

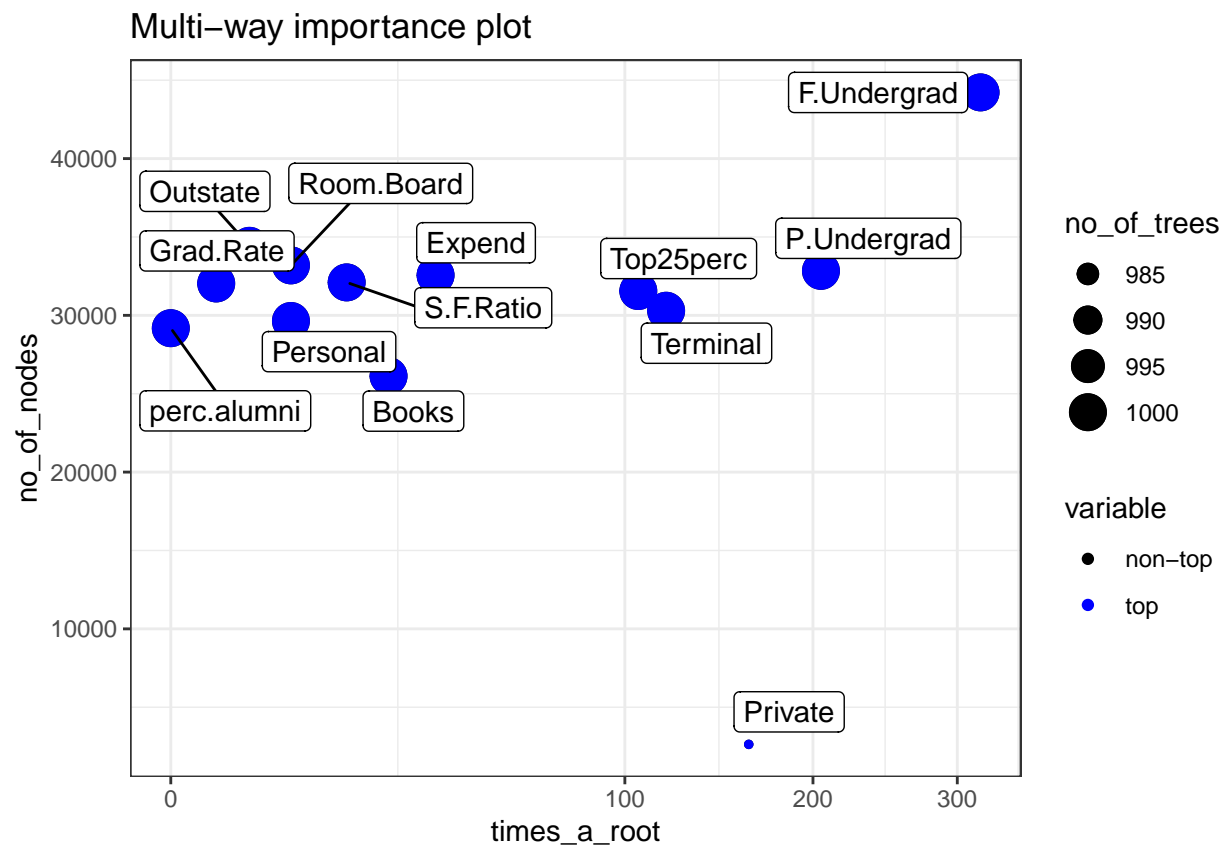
```
## [1] 1940.636
```

```
rf_rmse_test_set_1000
```

```
## [1] 1318.993
```

And now how the variables are used:

```
importance_random_forest_1000 <- measure_importance(random_forest_1000)
plot_multi_way_importance(importance_random_forest_1000, y_measure = "no_of_nodes", x_measure = "times_a_root")
```



Ten thousand trees model fitting

How said in the original paper, the random forest never overfits, so we can try to increase the number of trees to 10000:

```
random_forest_10000 <- randomForest(x=train_data, formula=Apps ~ ., y=College$Apps[train_index], data=train_data,
                                     ntree=10000)
random_forest_10000

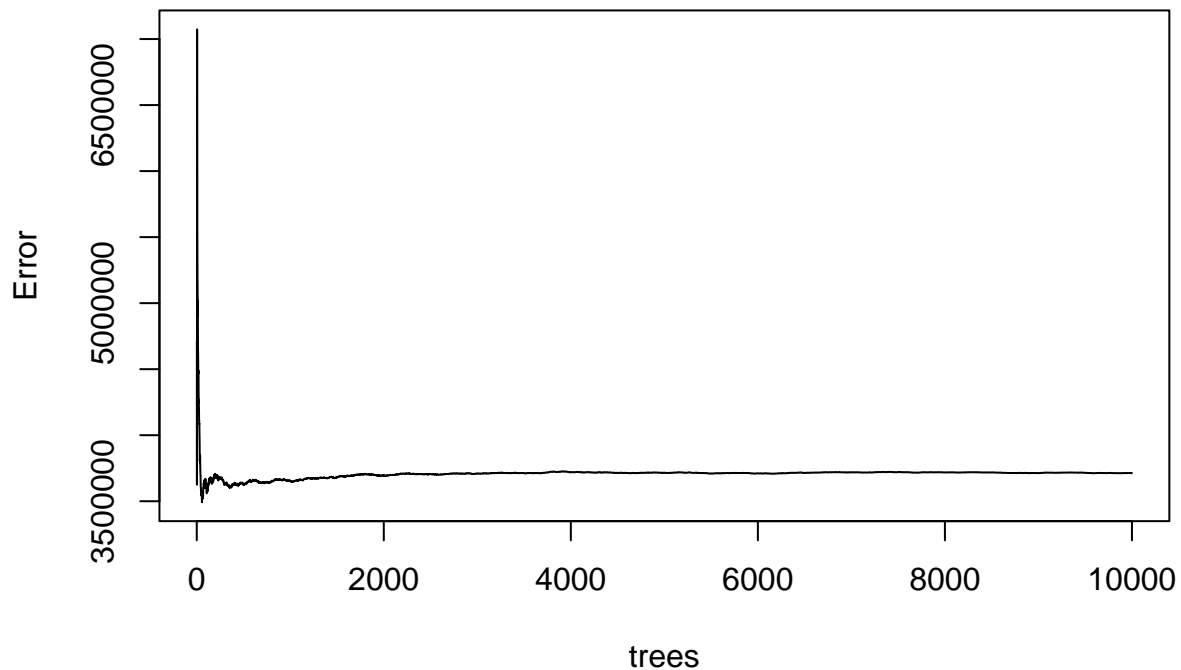
##
## Call:
## randomForest(x = train_data, y = College$Apps[train_index], xtest = test_data, ytest = College$Apps[test_index],
##              type = "regression", ntree = 10000,
##              mtry = 4,
##              nodesize = 1,
##              oob = TRUE,
##              importance = TRUE,
##              keep.forest = TRUE,
##              verbose = FALSE)
##              Type of random forest: regression
##              Number of trees: 10000
## No. of variables tried at each split: 4
##
##              Mean of squared residuals: 3713422
##              % Var explained: 76.91
##              Test set MSE: 1782928
##              % Var explained: 83.2
```

Ten thousand trees model evaluation

Now see how it performs:

```
plot(random_forest_10000, main = "Random forest MSE on train set vs number of trees")
```

Random forest MSE on train set vs number of trees



```
#RMSE train_set
rf_rmse_train_set_10000 <- sqrt(random_forest_10000$mse[length(random_forest_10000$mse)])

#RMSE test_set
rf_rmse_test_set_10000 <- sqrt(random_forest_10000$test$mse[length(random_forest_10000$test$mse)])

rf_rmse_train_set_10000
```

```
## [1] 1927.024
```

```
rf_rmse_test_set_10000
```

```
## [1] 1335.263
```


```
importance_random_forest_10000 <- measure_importance(random_forest_10000)

importance_random_forest_10000
```

```
##      variable mean_min_depth no_of_nodes mse_increase node_purity_increase
## 1      Books      3.774000     258049    -257220.2      214542716
## 2      Expend      2.611100     324372     1342183.3      606995063
## 3 F.Undergrad      1.126900     440661    14269744.0      3982724134
## 4   Grad.Rate      2.726600     322383     2969593.7      690952767
## 5   Outstate      2.657500     344772     1095401.2      501927836
```

## 6	P.Undergrad	2.208300	327868	3184373.1	1083674887
## 7	perc.alumni	3.880100	291053	407994.9	183954868
## 8	Personal	3.716300	298400	48999.6	221915378
## 9	Private	4.411572	25966	2326405.7	464815615
## 10	Room.Board	3.187900	331161	496765.5	327263613
## 11	S.F.Ratio	3.222300	321410	522618.1	340410084
## 12	Terminal	2.706400	304717	478133.4	498010888
## 13	Top25perc	2.269200	315546	994797.3	839793111
##	no_of_trees	times_a_root	p_value		
## 1	10000	227	1.000000e+00		
## 2	10000	346	0.000000e+00		
## 3	10000	3094	0.000000e+00		
## 4	10000	24	0.000000e+00		
## 5	10000	33	0.000000e+00		
## 6	10000	2184	0.000000e+00		
## 7	10000	5	1.000000e+00		
## 8	10000	78	9.999643e-01		
## 9	9798	1522	1.000000e+00		
## 10	10000	41	0.000000e+00		
## 11	10000	238	0.000000e+00		
## 12	10000	1114	5.739285e-16		
## 13	10000	1094	2.099544e-177		

```
plot_multi_way_importance(importance_random_forest_10000, y_measure = "no_of_nodes", x_measure = "times_a_root",
```



We can see that the random forest with 10000 trees perform slightly better the the previous one.


```
summary(glm)
```

Lets go back to the results for the glm model

```
##
## Call:
## glm(formula = College$Apps[train_index] ~ ., family = poisson,
##      data = train_data)
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  4.092305   0.008438  485.009 < 2e-16 ***
## Private      -0.184864   0.002964  -62.360 < 2e-16 ***
## Top25perc     0.177462   0.004820   36.818 < 2e-16 ***
## F.Undergrad   4.729292   0.007555  625.964 < 2e-16 ***
## P.Undergrad  -0.146204   0.007927  -18.443 < 2e-16 ***
## Outstate      0.492628   0.007140   68.997 < 2e-16 ***
## Room.Board    0.748332   0.006285  119.061 < 2e-16 ***
## Books          0.282714   0.010330   27.368 < 2e-16 ***
## Personal     -0.350038   0.008107  -43.179 < 2e-16 ***
## Terminal      0.019375   0.006914    2.802  0.00508 **
## S.F.Ratio     0.351358   0.009843   35.696 < 2e-16 ***
## perc.alumni  -0.049954   0.005894   -8.475 < 2e-16 ***
## Expend        0.570555   0.010381   54.959 < 2e-16 ***
## Grad.Rate     0.955605   0.006981  136.879 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for poisson family taken to be 1)
##
##      Null deviance: 2120079  on 618  degrees of freedom
## Residual deviance:  221113  on 605  degrees of freedom
## AIC: 226872
##
## Number of Fisher Scoring iterations: 4
```

There are a couple of takeaways here. First of all the the residual deviance is much lower than the null deviance, so the glm model explains a lot of variance. However the residual deviance od 221,113 is still quite large. Compared to the degrees of freedom this could be due to potential model misspecification, which could be due to non-linearity.

```
library(pdp)


names <- colnames(train_data)

names <- names[-1]

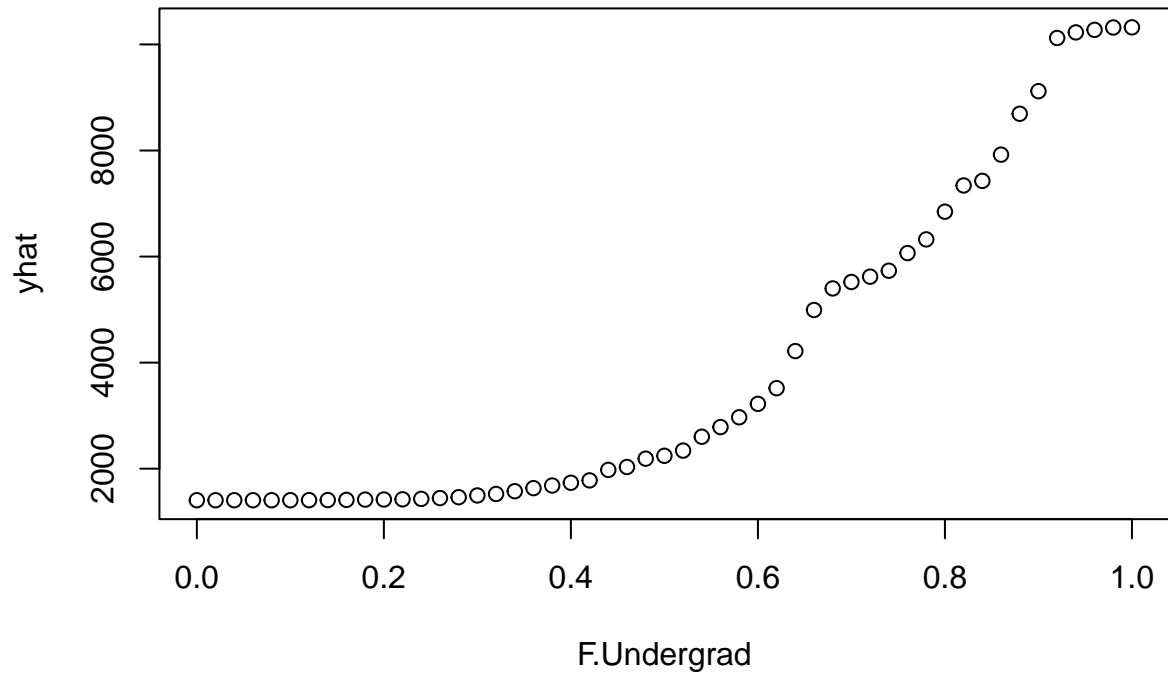
names
```

Now when we look at back at the multi way importance plot we see that the F.Undergrad and P.Undergrad are the most important features. Lets see the partial dependence plot for these two features.

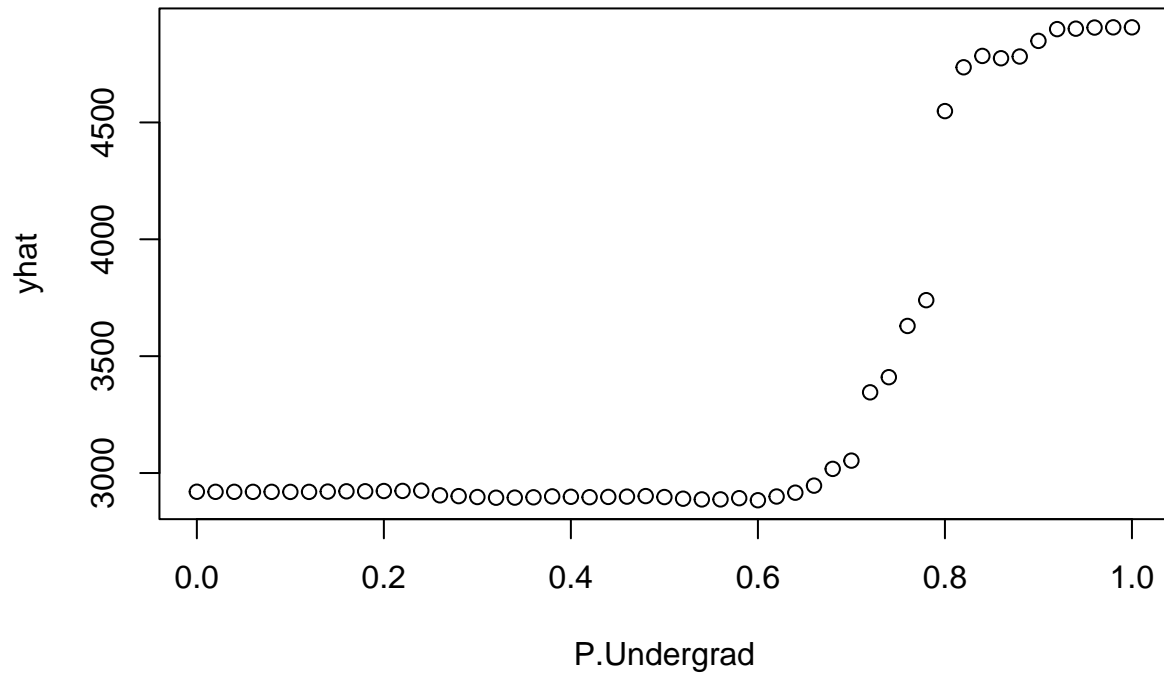
```
## [1] "Top25perc" "F.Undergrad" "P.Undergrad" "Outstate" "Room.Board"  
## [6] "Books" "Personal" "Terminal" "S.F.Ratio" "perc.alumni"  
## [11] "Expend" "Grad.Rate"  
  
# Partial dependence plot for F.Undergrad  
for (pred in names) {  
  pdp_obj <- partial(random_forest_1000, pred.var = pred, train = train_data)  
  plot(pdp_obj, main = paste("Partial Dependence for", pred))  
}
```



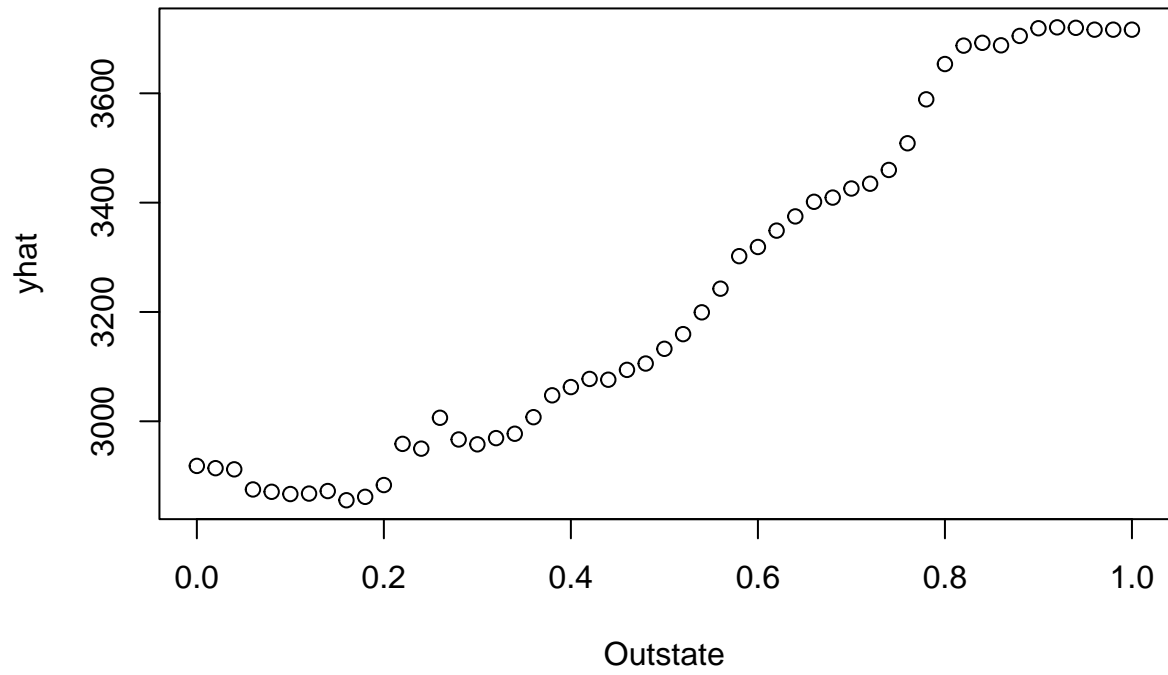
Partial Dependence for F.Undergrad




Partial Dependence for P.Undergrad



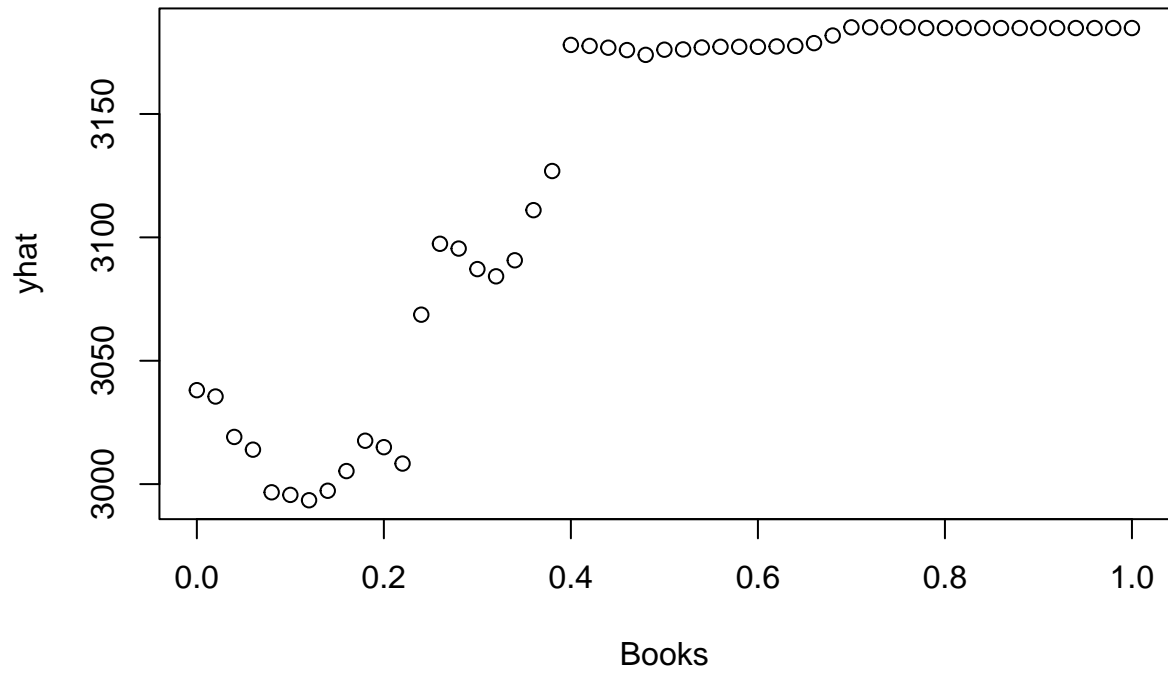
Partial Dependence for Outstate



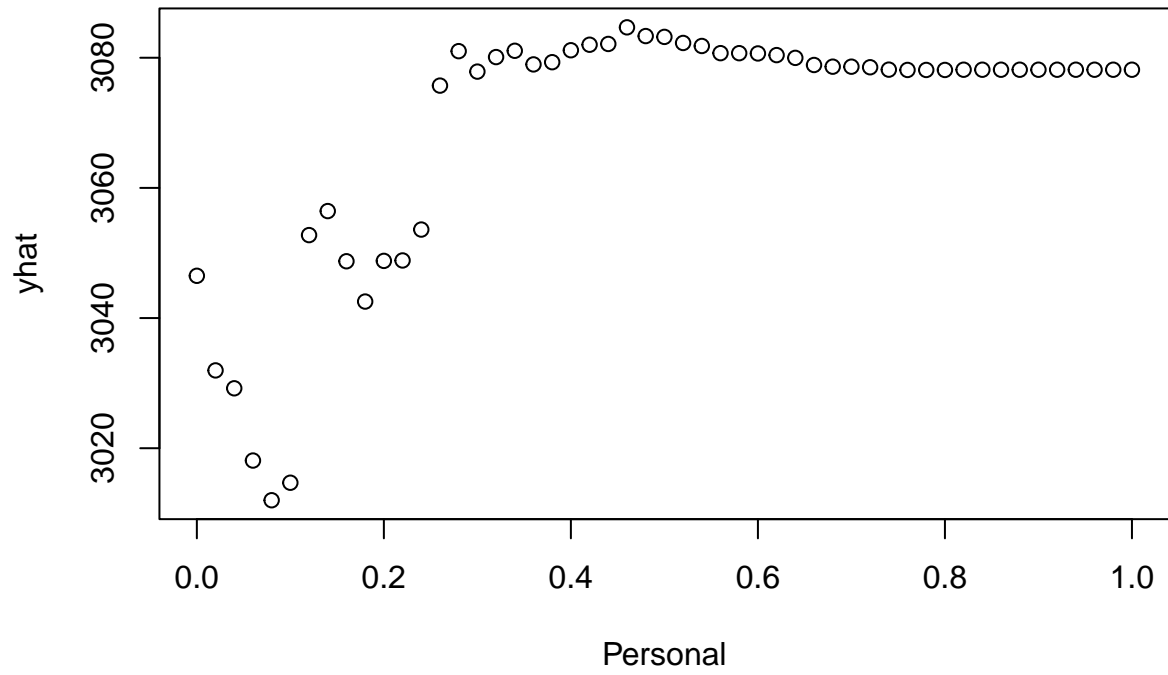
Partial Dependence for Room.Board



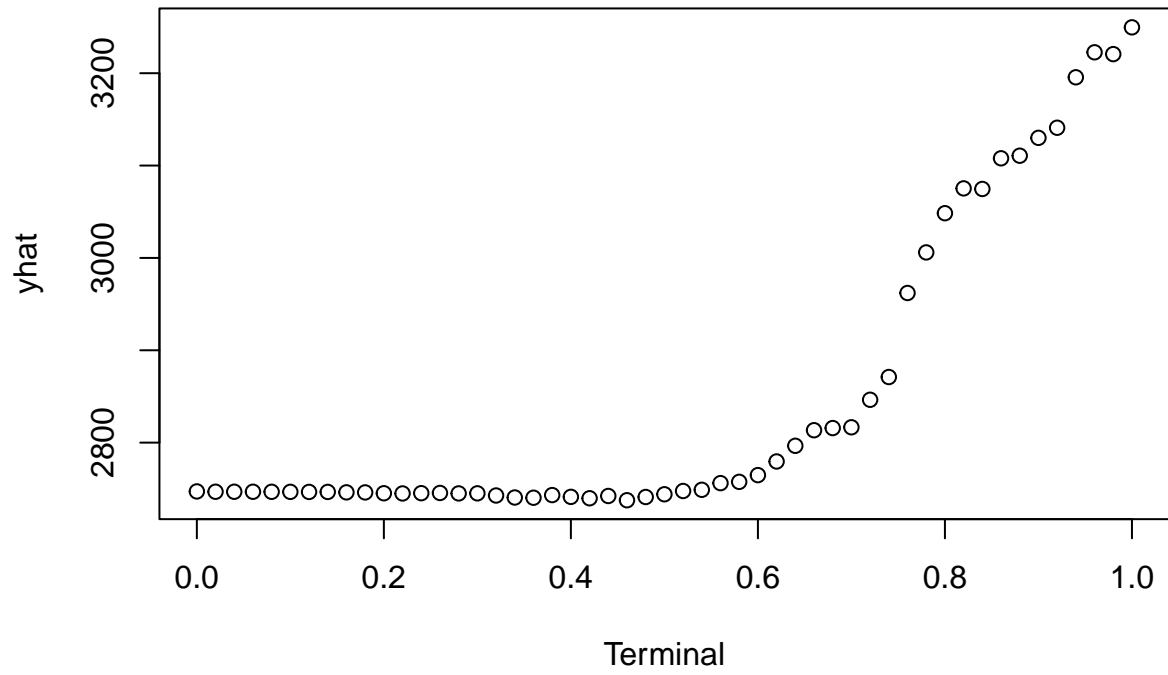
Partial Dependence for Books



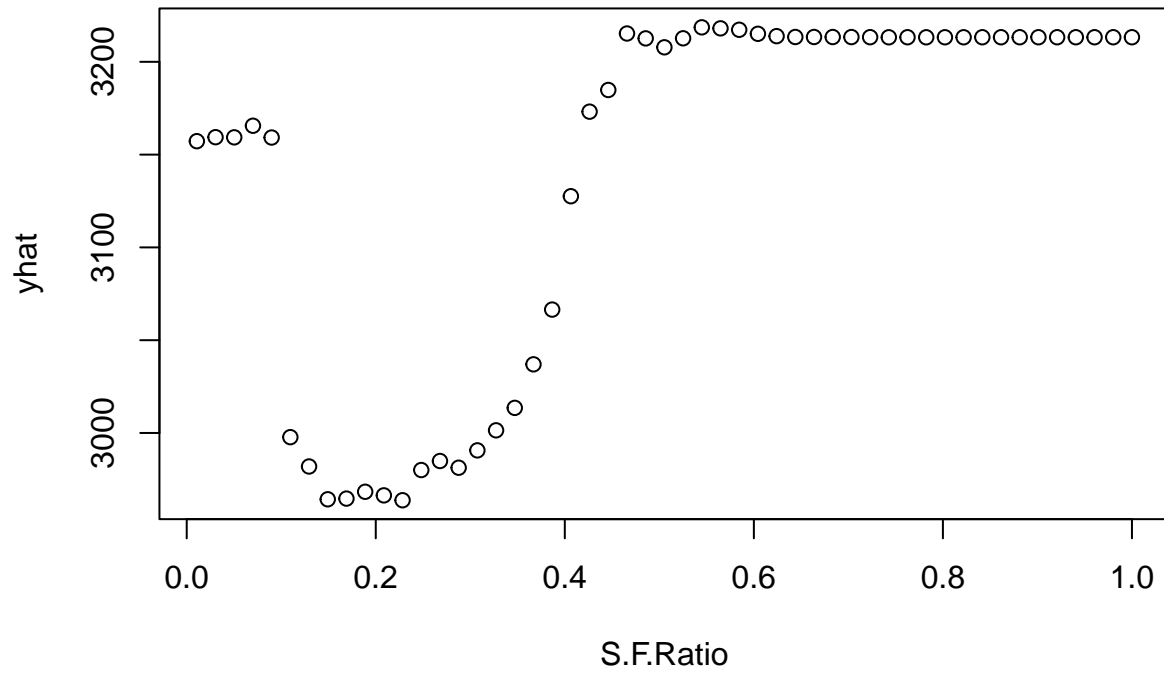
Partial Dependence for Personal



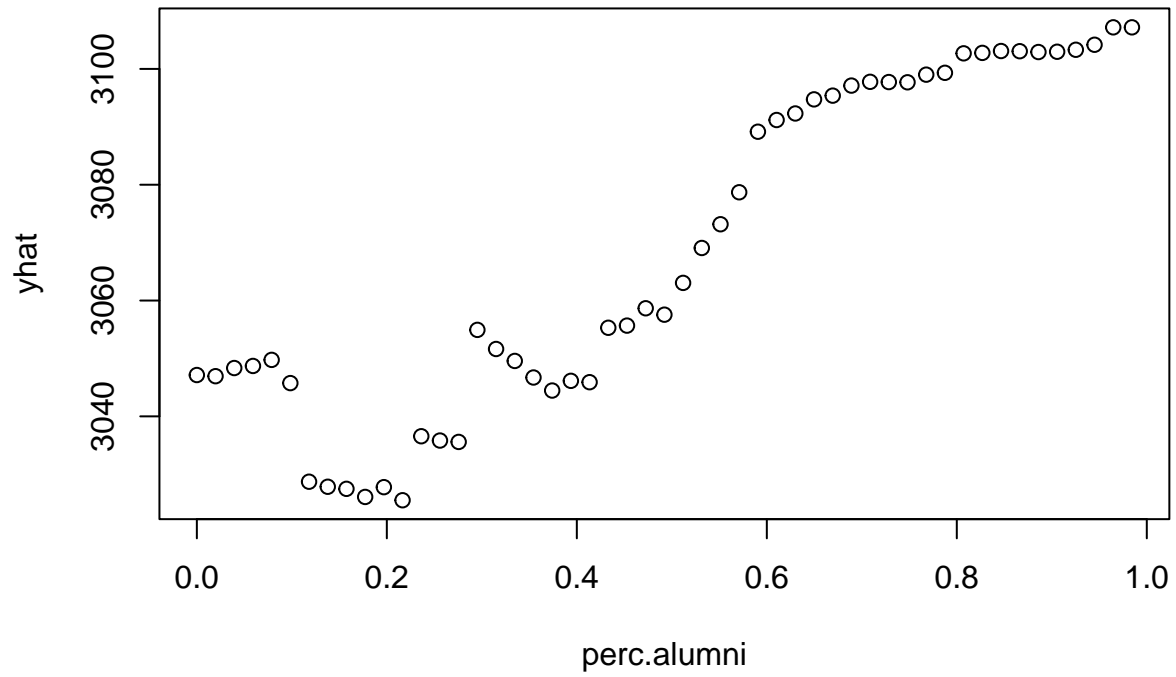
Partial Dependence for Terminal



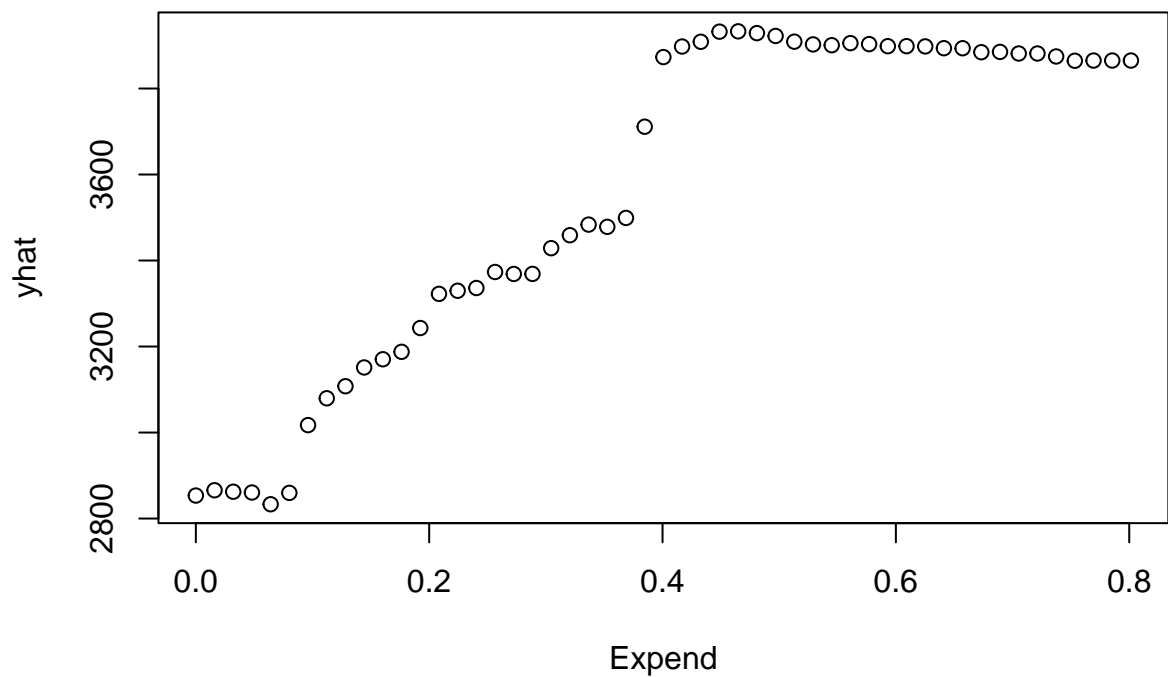
Partial Dependence for S.F.Ratio



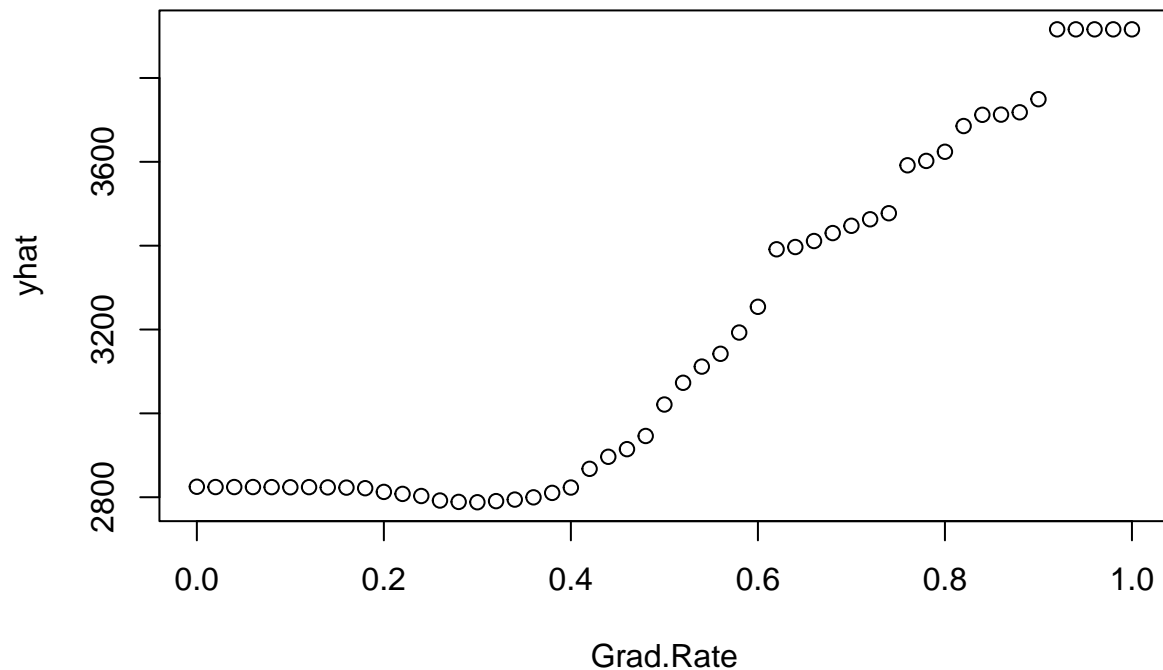
Partial Dependence for perc.alumni



Partial Dependence for Expend



Partial Dependence for Grad.Rate



Most of the features show us that they have a non-linear relationship with the response variable.

These two points regarding the GLM and the Random forest, suggest that a GAM could be a very strong alternative due to non-linearity between the features and the response variable.

```
library(mgcv)
```

Now let's fit a GAM on a training set

```
## Loading required package: nlme
```

```
##
```

```
## Attaching package: 'nlme'
```

```
## The following object is masked from 'package:dplyr':
```

```
##
```

```
## collapse
```

```
## This is mgcv 1.9-1. For overview type 'help("mgcv-package")'.
```

```
# Fit GAM model
gam_model <- gam(College$Apps[train_index] ~ F.Undergrad + P.Undergrad +
                Grad.Rate + Outstate +
                Room.Board + Expend +
                Books + Personal +
                S.F.Ratio + perc.alumni +
                Top25perc + Terminal,
                family = poisson, data = train_data)

summary(gam_model)
```

```
##
## Family: poisson
## Link function: log
##
## Formula:
## College$Apps[train_index] ~ F.Undergrad + P.Undergrad + Grad.Rate +
##      Outstate + Room.Board + Expend + Books + Personal + S.F.Ratio +
##      perc.alumni + Top25perc + Terminal
##
## Parametric coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  3.895692   0.007828  497.66  <2e-16 ***
## F.Undergrad  4.927114   0.006799  724.68  <2e-16 ***
## P.Undergrad -0.156095   0.007896  -19.77  <2e-16 ***
## Grad.Rate    0.957661   0.006934  138.10  <2e-16 ***
## Outstate     0.362536   0.006824   53.12  <2e-16 ***
## Room.Board   0.650742   0.006084  106.97  <2e-16 ***
## Expend       0.567961   0.010371   54.76  <2e-16 ***
## Books        0.219771   0.010278   21.38  <2e-16 ***
## Personal     -0.334710   0.008041  -41.62  <2e-16 ***
## S.F.Ratio    0.398119   0.009808   40.59  <2e-16 ***
## perc.alumni -0.124358   0.005744  -21.65  <2e-16 ***
## Top25perc    0.174991   0.004836   36.18  <2e-16 ***
## Terminal     0.140012   0.006645   21.07  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
##
## R-sq.(adj) =  0.821   Deviance explained = 89.4%
## UBRE =    362.6   Scale est. = 1           n = 619
```

```
gam_preds_test <- predict(gam_model, test_data, type = "response")

rmse_gam_test <- sqrt(mean((College$Apps[-train_index] - gam_preds_test)^2))

rmse_gam_test
```

And compare it to the glm

```
## [1] 1280.016
```

```
gam_preds_train <- predict(gam_model, train_data, type = "response")

rmse_gam_train <- sqrt(mean((College$Apps[train_index] - gam_preds_train)^2))

rmse_gam_train
```

```
## [1] 1678.38
```

We can also see a slight improvement in RMSE compared to the random forest models.

```
library(mgcv)
# Fit GAM model
gam_model_2 <- gam(College$Apps[train_index] ~ s(F.Undergrad, k=3) + s(P.Undergrad, k=3) +
                  Grad.Rate + Outstate +
                  Room.Board + s(Expend, k=3) +
                  s(Books, k=3) + s(Personal, k=3) +
                  s(S.F.Ratio, k=3) + s(perc.alumni, k=3) +
                  s(Top25perc, k=3) + s(Terminal, k=3),
                  family = poisson, data = train_data, select= TRUE)

summary(gam_model_2)
```

Now lets fine-tune the model a bit to see if we can get even better results

```
##
## Family: poisson
## Link function: log
##
## Formula:
## College$Apps[train_index] ~ s(F.Undergrad, k = 3) + s(P.Undergrad,
##   k = 3) + Grad.Rate + Outstate + Room.Board + s(Expend, k = 3) +
##   s(Books, k = 3) + s(Personal, k = 3) + s(S.F.Ratio, k = 3) +
##   s(perc.alumni, k = 3) + s(Top25perc, k = 3) + s(Terminal,
##   k = 3)
##
## Parametric coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  6.569381   0.004297 1528.76  <2e-16 ***
## Grad.Rate    1.037421   0.007355  141.04  <2e-16 ***
## Outstate     0.383205   0.007291   52.56  <2e-16 ***
## Room.Board   0.537761   0.006299   85.38  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Approximate significance of smooth terms:
##              edf Ref.df   Chi.sq p-value
## s(F.Undergrad) 1.9965     2 441509.2 <2e-16 ***
## s(P.Undergrad) 1.9821     2   992.8  <2e-16 ***
## s(Expend)       1.9914     2  1853.7  <2e-16 ***
```

```
## s(Books)          1.9957      2   1358.7 <2e-16 ***
## s(Personal)       1.9966      2   1988.3 <2e-16 ***
## s(S.F.Ratio)      1.9878      2   3599.0 <2e-16 ***
## s(perc.alumni)    1.9833      2    474.9 <2e-16 ***
## s(Top25perc)      1.9970      2   7404.2 <2e-16 ***
## s(Terminal)       0.9967      2    274.6 <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## R-sq.(adj) =  0.825   Deviance explained =  90%
## UBRE = 342.49   Scale est. = 1           n = 619
```

```
gam_preds_test_2 <- predict(gam_model_2, test_data, type = "response")

rmse_gam_test_2 <- sqrt(mean((College$Apps[-train_index] - gam_preds_test_2)^2))

rmse_gam_test_2
```

And lets see the RMSE

```
## [1] 1308.551
```

```
gam_preds_train_2 <- predict(gam_model_2, train_data, type = "response")

rmse_gam_train_2 <- sqrt(mean((College$Apps[train_index] - gam_preds_train_2)^2))

rmse_gam_train_2
```

```
## [1] 1649.383
```

```
AIC(gam_model, gam_model_2)
```

Comparing the models

```
##              df      AIC
## gam_model    13.00000 230796.0
## gam_model_2  20.92708 218353.6
```

Comparing predictions of the GAM vs Random forest

```
plt_num <- length(College$Apps[-train_index])-1

plot(0:plt_num, gam_preds_test, col = "blue", pch = 16,
     main = "GAM vs. Random Forest Predictions",
```



```

xlab = "School Index", ylab = "Applications")

points(0:plt_num, predict(random_forest_1000, test_data), col = "red", pch = 16)
points(0:plt_num, College$Apps[-train_index], col = "black", pch = 16)

for (i in 0:plt_num) {
  abline(v = i, col = "gray", lwd = 0.5, lty = 2) # Dashed thin vertical lines
}

legend("topleft",inset=c(0,-0.1), legend = c("Actual Value", "Random Forest","GAM"),
      col = c("black", "red", "blue"), pch = c(16, 16, 16), bty = "n", horiz=TRUE,xpd=TRUE)

```

